

PARAREAL ALGORITHM IMPLEMENTATION AND SIMULATION IN JULIA

Tyler M. Masthay[†] and Saverio Perugini[‡]

[†]Institute for Computational and Engineering Sciences, University of Texas Austin, tyler@ices.utexas.edu

[‡]Department of Computer Science, University of Dayton, Ohio 45469-2160, saverio@udayton.edu

August 16, 2018

SUMMARY

We present an implementation of the parareal algorithm—an integration technique to solve differential equations in parallel—first proposed in 2001 by Lions, Maday, and Turinici [3]—in the Julia programming language [4] for a fully general, first-order, initial-value problem. We also provide a graphical simulation of the parareal algorithm intended to be used in a numerical analysis course to both introduce the parareal algorithm to students and aid them in an investigation of the types of curves for which the parareal algorithm might be practical. Our implementation of the parareal algorithm accepts both coarse and fine integrators as functional arguments. We provide implementations of Euler’s method and another Runge-Kutta integration technique as the integrators. The final two functions in the source code file `parareal.jl` are functions implementing Euler’s method and another Runge-Kutta integration technique that can be used as examples to be passed as first-class functions as coarse or fine integration techniques to the `parareal` or `simulate` functions. A Git repository of both the implementation and graphical simulation is available in GitHub at <https://github.com/sperugin/Parareal-Implementation-and-Simulation-in-Julia.git>. All of the graphical plots are generated with the Julia Plots package available at <https://juliaplots.github.io/>. A video describing this application of Julia is available on YouTube at <https://www.youtube.com/watch?v=MtgbeL06ZM4>. The purpose of this software is pedagogical: as a simulation to introduce students to the parareal algorithm and the concept of concurrency, and as a tool for (graphically) investigating the performance of the algorithm.

Keywords: Concurrent programming, *Euler’s method*, Julia, *Runge-Kutta* methods, parareal algorithm, ordinary differential equations.




Figure 1: Right endpoint error.

THE PARAREAL ALGORITHM

The parareal algorithm is designed to perform parallel-in-time integration for a first-order initial-value problem. The algorithm involves two integration techniques, often known as the ‘coarse’ integrator and the ‘fine’ integrator. For the algorithm to be effective, the coarse integrator must be of substantially lower computational cost than the fine integrator. The reason will become apparent later in this section. Consider the differential equation (1) given by

$$y'(t) = f(t, y(t)) \quad t \in [a, b] \quad (1)$$

with its associated initial-value problem (2)

$$y(t^*) = y^* \quad t^* \in [a, b]. \quad (2)$$

For simplicity, let us assume $t^* = a$, so that the solution only extends rightward. To obtain an approximate solution to equation (1) satisfying the initial condition (2), we partition our domain into $[t_0 = a, \dots, t_N = b]$ with uniform step size Δ . We now precisely define an ‘integrator’ as a function from $(0, \infty) \times \mathbb{R}^2 \times \mathcal{R}$ to \mathbb{R} where \mathcal{R} is the set of all Riemann integrable functions. For example, the integrator I given by

$$I(\delta, x_0, y_0, g) = y_0 + g(x_0, y_0)\delta$$

is the integrator corresponding to Euler’s method with step size δ . Let \mathcal{C} and \mathcal{F} be the coarse and fine integrators, respectively. Define

$$\begin{aligned} y_{0,1} &= y(t_0) = y^* \\ y_{n+1,1} &= y(t_{n+1}) = \mathcal{C}(\Delta, t_n, y_{n,1}, f). \end{aligned}$$

Since $y_{n+1,1}$ depends on $y_{n,1}$, this algorithm is inherently sequential. Partition $[t_n, t_{n+1}]$ into $\{t_n^0 = t_n, \dots, t_n^m, \dots, t_n^M = t_{n+1}\}$ with uniform step size $\delta < \Delta$. Define

$$\begin{aligned} z_{n,1}^0 &= y(t_n^0) = y_{n,1} \\ z_{n,1}^{m+1} &= y(t_n^{m+1}) = \mathcal{F}(\delta, t_n^m, z_{n,1}^m, f). \end{aligned}$$

This yields an approximate solution $\{z_{n,1}^0, \dots, z_{n,1}^M\}$ to (1) over $[t_n, t_{n+1}]$ with initial conditions

$$y(t_n) = y_{n,1}.$$

Since $z_{n_1,1}^{m_1}$ does not depend on $z_{n_2,1}^{m_2}$ for $n_1 \neq n_2$, we can compute these approximations in parallel. After the last subproblem is solved, we simply combine the solutions on each subdomain to obtain a solution over the whole interval. However, our values $\{y_{1,1}, \dots, y_{n,1}\}$ are relatively inaccurate. The vertical spikes in the orange graph separating the coarse and fine predictions in Figure 1 illustrate this error. However, $z_{n-1,1}^M$ is a better approximation for $\phi(t_n)$ where ϕ is the exact solution to the differential equation. We use this to obtain a better set of points $\{y_{n,2}\}$ for the coarse approximation. We do this by first defining $w_{n,1} = y_{n,1}$ and then defining

$$\begin{aligned} w_{1,2} &= y_{1,1} = y_{1,2} = y^* \\ w_{n,2} &= \mathcal{C}(\Delta, t_{n-1}, y_{n-1,2}, f) \\ y_{n,2} &= (w_{n,2} - w_{n,1}) + z_{n-1,1}^M. \end{aligned}$$

Thus, $w_{n+1,2}$ serves as a new prediction given a more accurate previous prediction from $y_{n,2}$ since $z_{n-1,1}^M$ has now been taken into account in calculating $y_{n,2}$. In general, we continue evaluating so that for $k > 1$, we have

$$\begin{aligned} w_{1,k} &= y_{1,k} = y^* \\ w_{n,k} &= \mathcal{C}(\Delta, t_{n-1}, y_{n-1,k-1}, f) \\ y_{n,k} &= (w_{n,k} - w_{n,k-1}) + z_{n-1,k-1}^M. \end{aligned}$$

Note that since $y_{n,k}$ is dependent on $w_{n,k}$, this step must be done sequentially. As k increases, $w_{n,k} - w_{n,k-1} \rightarrow 0$, which means that $y_{n,k}$ converges to the value that the fine integrator would predict if fine integration were simply done sequentially. Thus, each k denotes fine integration over the whole interval. This means that the total computation performed is much greater than if fine integration were performed sequentially. However, the time efficiency of each iteration has the potential to be improved through concurrency. Since fine integration is more computationally intensive, this improvement in the run-time efficiency may compensate for the cumulative computation performed.




Figure 2: Slow parareal example. (left) Solution after first iteration with Euler's method. (right) Solution after ninth iteration with Euler's method.

Let K be the total number of iterations necessary to achieve a desired accuracy of solution and P be the number of subintervals into which we divide according to the coarse integrator. If $K = 1$, then we achieve perfect parallel efficiency. If $K = P$, then we likely slowed the computation down. The parareal algorithm is guaranteed to converge to the solution given by the sequential fine integrator within P iterations. For a more complete treatment of this convergence analysis, we refer the reader to [2]. For fully general pseudocode, we refer the reader to [1, 5].

PARAREAL ALGORITHM IMPLEMENTATION IN JULIA

The `@async` macro within the loop causes the program to evaluate the first expression to its right as a concurrent task (i.e., the `for` loop assigning values to `sub`). The `@sync` macro causes the main program thread to wait until all tasks (spawned in the the first expression to its right with an `@async` or `@parallel` macro) complete. Once all concurrent tasks are complete, execution of the program proceeds sequentially. Given the semantics of these macros, the program correctly perform concurrent integration. The sequential and parallel versions of this implementation have no significant differences in run-time efficiency. However, if a `sleep` statement is placed in the argument of `fineIntegrator`, the parallel version runs much faster. This demonstrates that use of those two macros leads to concurrent program execution.

GRAPHICAL ALGORITHM SIMULATION

The function `simulate` creates a graphical simulator of the parareal algorithm and can be used to introduce the parareal algorithm to students in a numerical analysis course. The first line gets the sequential solution from the fine integrator (the ‘ideal’ solution) and the second line gets the history of the computations that took place during the parareal execution. The main loop over the variable k then displays the inner workings of the algorithm. The ideal solution is plotted, with a scatter plot of the points obtained from the coarse integrator. To simulate the parallel nature of the algorithm, random progress is made on randomly selected subdomains. Thus, the plot dynamically makes partial progress on different subdomains until all subdomains are finished with the fine integration. After this, the plots are connected into the current iteration’s approximation. During the next iteration, the previous guesses from the coarse integrator are displayed in red and the new guesses from the coarse integrator are displayed in green. As k increases, these guesses converge to the ideal solution.

In addition to the use of this function for pedagogical purposes, it can be used to investigate the types of curves for which the parareal algorithm might be practical. For instance, consider the differential equation

$$y'(x) = \sin(xy), \quad x \in [-20, 20]$$

with $y(-20) = 10$, $\Delta = 4$ (10 points), and $\delta = 0.008$ (500 points). Figure 2 shows the first and ninth iterations. The ninth iteration’s large error on the right end of the interval shows that this is an example where parareal convergence is slow. This is as inefficient as possible, needing as many iterations as subdomains in order for the solution to converge. However, the simulation also shows that if $f(x, y) = \sin(x)e^x$, then the solution converges after just one iteration. These two examples show that the algorithm’s efficiency can be highly dependent on the integrand.

REFERENCES

- [1] E. Aubanel. Scheduling of tasks in the parareal algorithm. *Parallel Computing*, 37(3):172–182, 2011.
- [2] M.J. Gander and S. Vandewalle. Analysis of the parareal time-parallel time-integration method. *SIAM Journal on Scientific Computing*, 29(2):556–578, 2007.

- [3] J.-L. Lions, Y. Maday, and G. Turinici. A “parareal” in time discretization of pde’s. *Comptes Rendus de l’Académie des Sciences - Series I - Mathematics*, 332:661–668, 2001.
- [4] J. Moffit and B.A. Tate. Julia. In B.A. Tate, F. Daoud, I. Dees, and J. Moffit, editors, *Seven more languages in seven weeks: Languages that are shaping the future*, chapter 5, pages 171–207. Pragmatic Bookshelf, Dallas, TX, 2014.
- [5] A.S. Nielsen. Feasibility study of the parareal algorithm. Master’s thesis, Technical University of Denmark, 2012.